


ERNIE (Enhanced Representation
through Knowledge Integration)

Big Bird: Transformers for
Longer Sequences

BERT (Bidirectional
Encoder Representations
from Transformers)

ELMo
(Embeddings from
Language Models)



STAYREAL



OSCAR: Explai(o)nable AI in Trash Classification(r)

Samara Goltz and
Bryn Archambault

What Is Our Project?

Trash Classification Using Convolutional Neural Networks

Project Category: Computer Vision

Yujie He
ICME
Stanford University
yujiehe@stanford.edu

Qinyue Gu
ICME
Stanford University
gqy94@stanford.edu

Maguo Shi
ICME
Stanford University
smgyl@stanford.edu

Abstract

As the waste problem becomes increasingly eminent across the globe, we aim to provide an automated waste sorting tool to make it easier for residents to classify trash. Our project used TrashNet [2] as our dataset, and classified recyclables or trash into six categories. To achieve our objective, we focused on Convolutional Neural Networks (CNN), and explored several well-known architectures at early stages. We ended up with modified AlexNet by taking two layers out, and experimented different techniques based on this model architecture, including dropout, data augmentation and learning rate decay. We also experimented two classifiers, Softmax and Support Vector Machine (SVM), as the last layer of our model structure. The highest test accuracy we achieved was 79.94% with the model using partial data augmentation and SVM classifier.

Problem Statement: What aspects of an image does AlexNet focus on when classifying an object within the TrashNet dataset? Additionally, how do different aspects of the model impact this?

Our Baseline Model

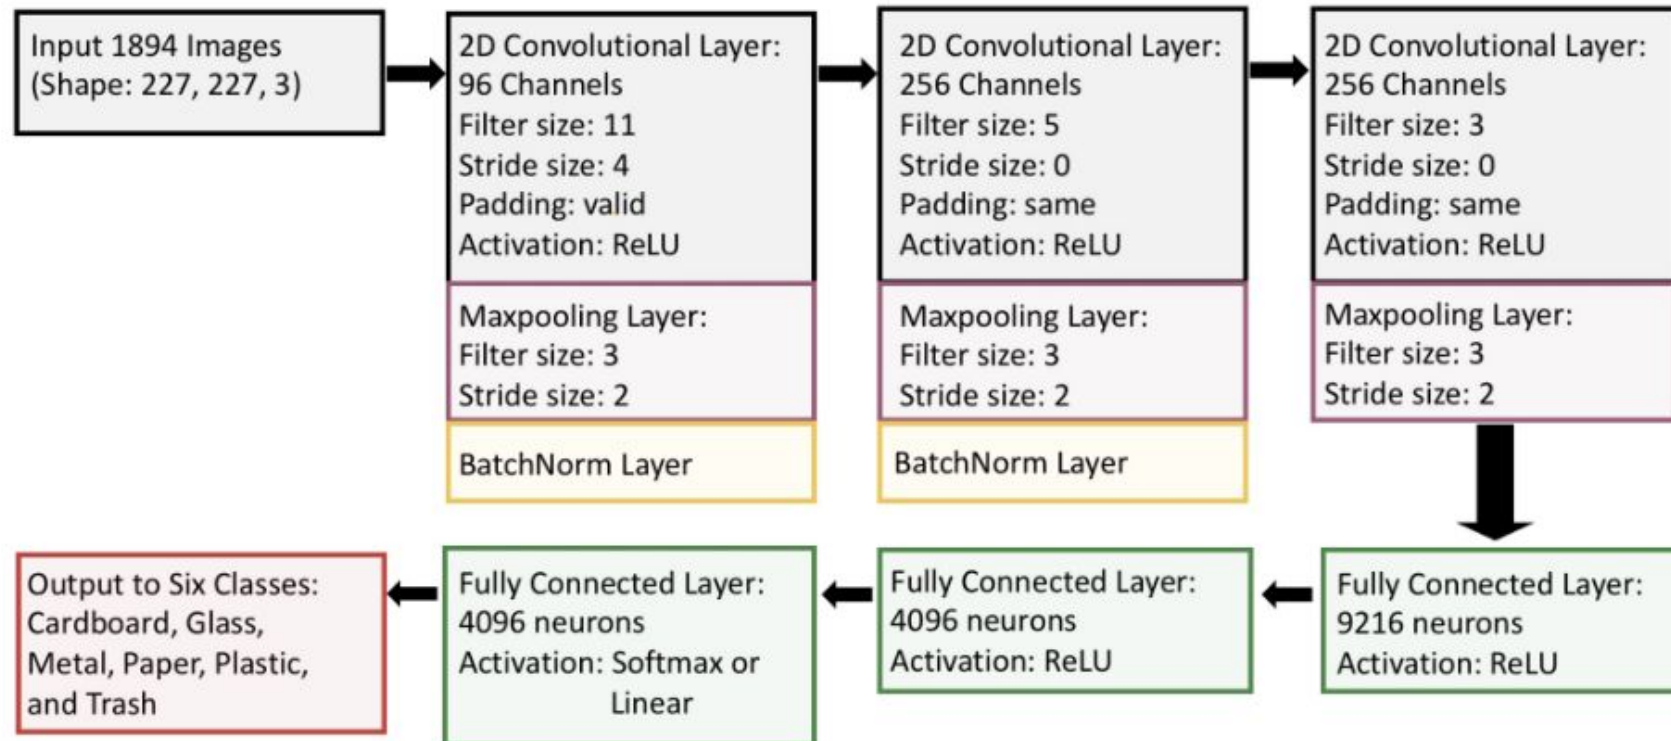


Figure 2: Basic Model Architecture

Comparison Plan



Using local explainable AI, perform ablation study using AlexNet Milestone #2 as the Baseline.

Focus on:

- Dropout
- Normalization
- ReLu
- CNN layers
- Max pooling - Resource Exhaustion error, requires too much memory
- Wanted to compare horizontally with AlexNet vs ResNet & VGG - Was going to take 16 & 28 hours to run

Compare:

- Accuracy
- Loss
- One explainable AI output image from each trash category

Our Inputs

TrashNet is a popular dataset for trash classification (6 features; 2,527 images):

- size (227 by 227)
- Plain background
- Trash is used, not perfect



Dataset Troubleshooting

```
import os
import numpy as np
import tensorflow as tf

# Path to the "cardboard" subfolder
cardboard_folder_path = '/content/drive/MyDrive/Senior Year/Machine Learning/Final Project/dataset-resized/cardboard'

# Check if the folder exists
if os.path.exists(cardboard_folder_path) and os.path.isdir(cardboard_folder_path):
    cardboard = []
    for file_name in os.listdir(cardboard_folder_path):
        image_path = os.path.join(cardboard_folder_path, file_name)
        temp = tf.keras.preprocessing.image.load_img(
            path=image_path,
            color_mode='rgb',
            target_size=(227, 227)
        )
        X = np.array(temp)
        cardboard.append(X)

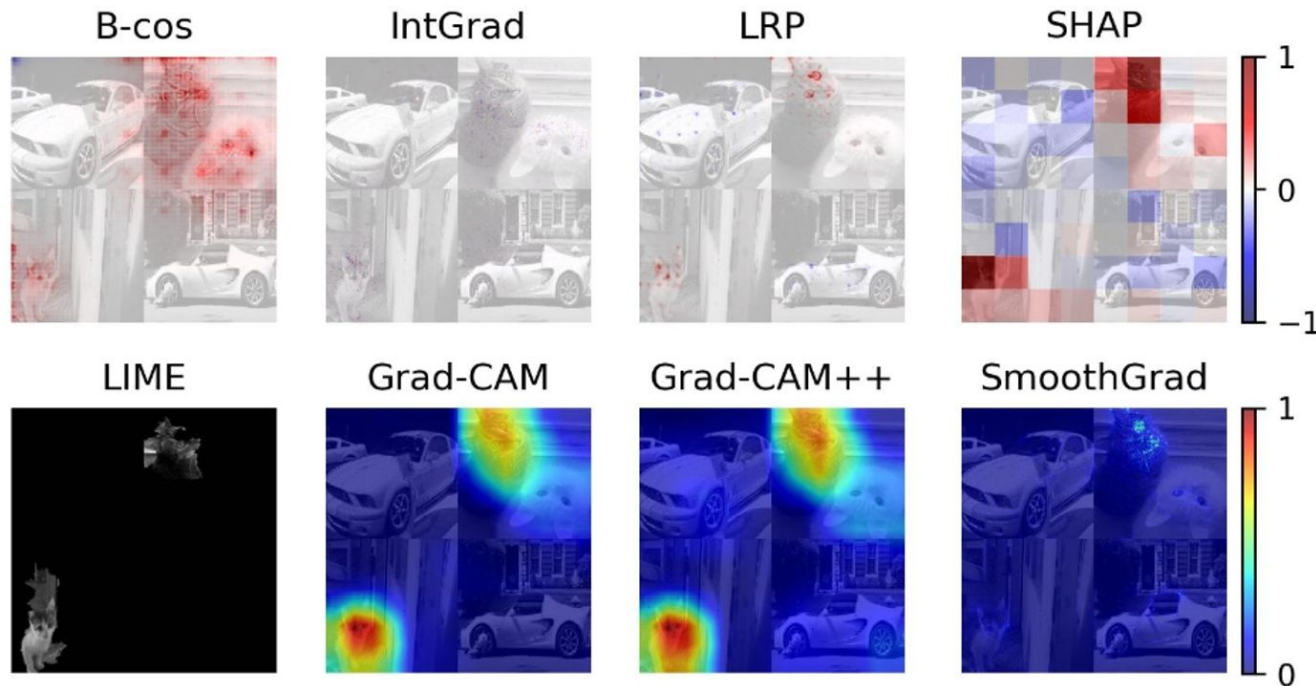
    cardboard = np.array(cardboard)
    cardboard = np.take(cardboard, np.random.permutation(cardboard.shape[0]), axis=0)
    print(cardboard.shape)
else:
    #print("The 'cardboard' folder does not exist or the path is incorrect.")
    cardboard_folder_path = '/content/drive/MyDrive/Junior Year/Machine Learning/Final Project/dataset-resized/cardboard'
    if os.path.exists(cardboard_folder_path) and os.path.isdir(cardboard_folder_path):
        cardboard = []
        for file_name in os.listdir(cardboard_folder_path):
            image_path = os.path.join(cardboard_folder_path, file_name)
            temp = tf.keras.preprocessing.image.load_img(
                path=image_path,
                color_mode='rgb',
                target_size=(227, 227)
            )
            X = np.array(temp)
            cardboard.append(X)

        cardboard = np.array(cardboard)
        cardboard = np.take(cardboard, np.random.permutation(cardboard.shape[0]), axis=0)
        print(cardboard.shape)
    else:
        print("The 'cardboard' folder does not exist or the path is incorrect.")
```


What is Explainable AI

What is it?

Explainable AI allows us to understand the inner workings of the model by opening up the “black box” to see what the model is focusing on



Benefits

- Trust in the model, important when considering money, health, safety, etc
- Sorting trash correctly is beneficial to the environment and is something simple everyone can do

Different Types

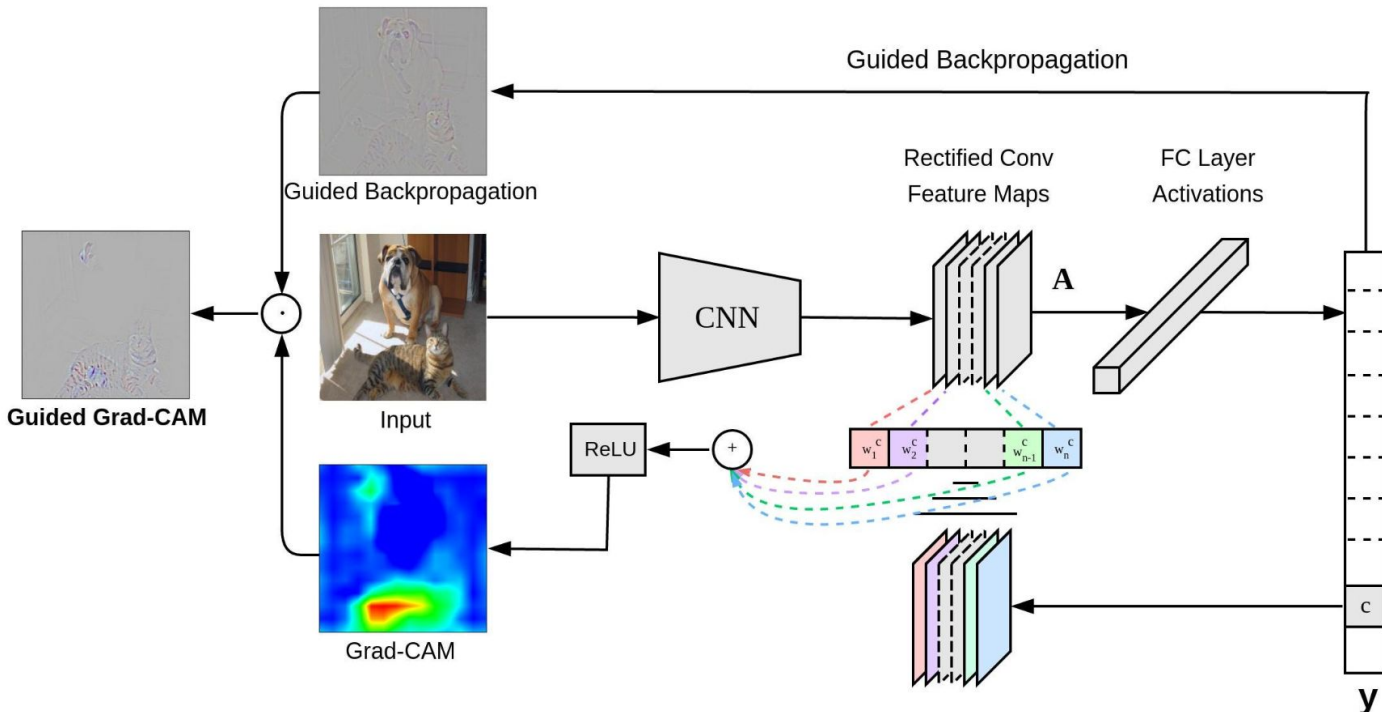
- Saliency Maps
- LIME
- LRP
- Decision Trees (Forests) - Global

What is Grad-CAM

What is it?

Inputs: A trained CNN model (AlexNet), an input image, and the target class index.

Outputs: A heatmap showing which parts of the image influenced the prediction



How Does It Work?

1. Focuses on the last convolutional layer
2. Passes image through network, assigning it to a class (trash, glass, etc)
3. The **gradient** of the score for a specific class shows how much each neuron contributed to that score
4. The last CNN outputs a **feature map** that uses gradients to highlight the influence of a feature.
 - a. Grad-CAM calculates the importance of each map, and weights them according to the gradients.
5. The weighted maps are combined into an **activation map**, and then normalized into a **heatmap**
6. Grad-CAM combines original image with heatmap

What is Grad-CAM

What is it doing?

1. Retrieves the output of the last conv. layer
2. Computes the gradient
3. Averages gradients spatially to get “importance weights” (pooled gradients) for each channel
4. Weighs the layer’s outputs by these gradients to create a heatmap
5. Rescales and normalizes heatmap for visualization

```
[ ] def grad_cam(model, image, class_idx):  
    """Generates a heat map using Grad-CAM."""  
    print(f"Inside grad cam, image shape: {image.shape}")  
    # Expand dimensions of the image to match input shape  
    image = np.expand_dims(image, axis=0) # Ensure batch dimension is included  
    # Get the gradient model  
    grad_model = tf.keras.Model(  
        inputs=[model.inputs], outputs=[model.get_layer("conv5").output, model.output])  
  
    with tf.GradientTape() as tape:  
        last_conv_layer = model.get_layer('conv5') # Ensure this matches your model's last conv layer  
        iterate = tf.keras.models.Model([model.input], [last_conv_layer.output, model.output])  
  
        # Use the already expanded image  
        conv_outputs, preds = iterate(image) # Image is now (1, 227, 227, 3)  
        loss = preds[:, class_idx]  
  
        grads = tape.gradient(loss, conv_outputs)[0]  
        pooled_grads = tf.reduce_mean(grads, axis=(0, 1)) # Average over the spatial dimensions  
        # Convert the conv_outputs to Numpy array  
        heatmap = conv_outputs[0].numpy() # Convert to Numpy array for manipulation  
  
        print(f"Shape of pooled_grads: {pooled_grads.shape}") # Debugging line  
        print(f"Shape of heatmap before scaling: {heatmap.shape}") # Debugging line  
  
        # Ensure pooled_grads is a numpy array  
        pooled_grads = pooled_grads.numpy() # Convert to Numpy for manipulation  
        pooled_grads = np.expand_dims(pooled_grads, axis=(0, 1)) # Shape becomes (1, 1, 256)  
  
        # Check shapes before multiplication  
        print(f"Shape of pooled_grads after expansion: {pooled_grads.shape}") # Debugging line  
        print(f"Shape of heatmap before applying pooled_grads: {heatmap.shape}") # Debugging line  
  
        # Multiply each channel in the heatmap with the corresponding pooled gradient  
        # Ensure correct dimensions for multiplication  
        heatmap = np.tensordot(heatmap, pooled_grads, axes=(2, 2)) # Shape will be (13, 13)  
        heatmap = np.maximum(heatmap, 0) # Ensure non-negative values  
        heatmap /= np.max(heatmap) # Normalize to [0, 1]  
  
        # Squeeze the heatmap  
        heatmap = np.squeeze(heatmap) # Remove dimensions of size 1  
        # Check heatmap shape before resizing  
        print(f"Heatmap shape before resizing: {heatmap.shape}")  
        # Resize heatmap to match the input image dimensions  
        if heatmap.ndim == 2: # Ensure the heatmap is 2D  
            heatmap = cv2.resize(heatmap, (image.shape[2], image.shape[3])) # Adjust to your image shape  
        else:  
            raise ValueError("Heatmap does not have the expected shape for resizing.")  
    return heatmap
```

What is Grad-CAM

What is it doing?

Inputs: The original image and the heatmap

Outputs: Visualization of the heatmap over the original image

1. Converts heatmap to RGB color map
2. Blends the heatmap with the original image to create a superimposed image
3. Displays the superimposed image using matplotlib

```
def display_grad_cam(image, heatmap):  
    # Resize heatmap to match the original image size  
    heatmap = np.uint8(255 * heatmap) # Scale to 0-255  
    heatmap = cv2.applyColorMap(heatmap, cv2.COLORMAP_JET) # Apply color map  
  
    # Overlay the heatmap on the original image  
    superimposed_img = heatmap * 0.4 + image # Adjust alpha as needed  
    plt.imshow(superimposed_img / 255.0) # Display as float image  
    plt.axis('off')  
    plt.show()  
  
# Run Grad-CAM on the selected images  
print(f"Number of selected images: {len(selected_images)}") # Debug statement  
  
for idx, selected_image in enumerate(selected_images):  
    print(f"Selected image {idx} shape: {selected_image.shape}") # Debug statement  
  
    # Check if the loop is entering the body  
    if selected_image is None:  
        print(f"Image {idx} is None.") # Check for None  
        continue # Skip if None  
  
    # Ensure we expand the dimensions correctly for the model  
    expanded_image = np.expand_dims(selected_image, axis=0)  
    print(f"After expansion, expanded image shape: {expanded_image.shape}")  
  
    # Predict class using the model  
    class_idx = np.argmax(model.predict(expanded_image))  
    print(f"Image {idx}: Predicted class index: {class_idx}")  
  
    # Call Grad-CAM  
    print(f"Calling grad_cam for image {idx} with class index {class_idx}")  
    heatmap = grad_cam(model, np.squeeze(selected_image), class_idx) # Remove extra dimensions  
    display_grad_cam(selected_image, heatmap)
```


Our Baseline Model

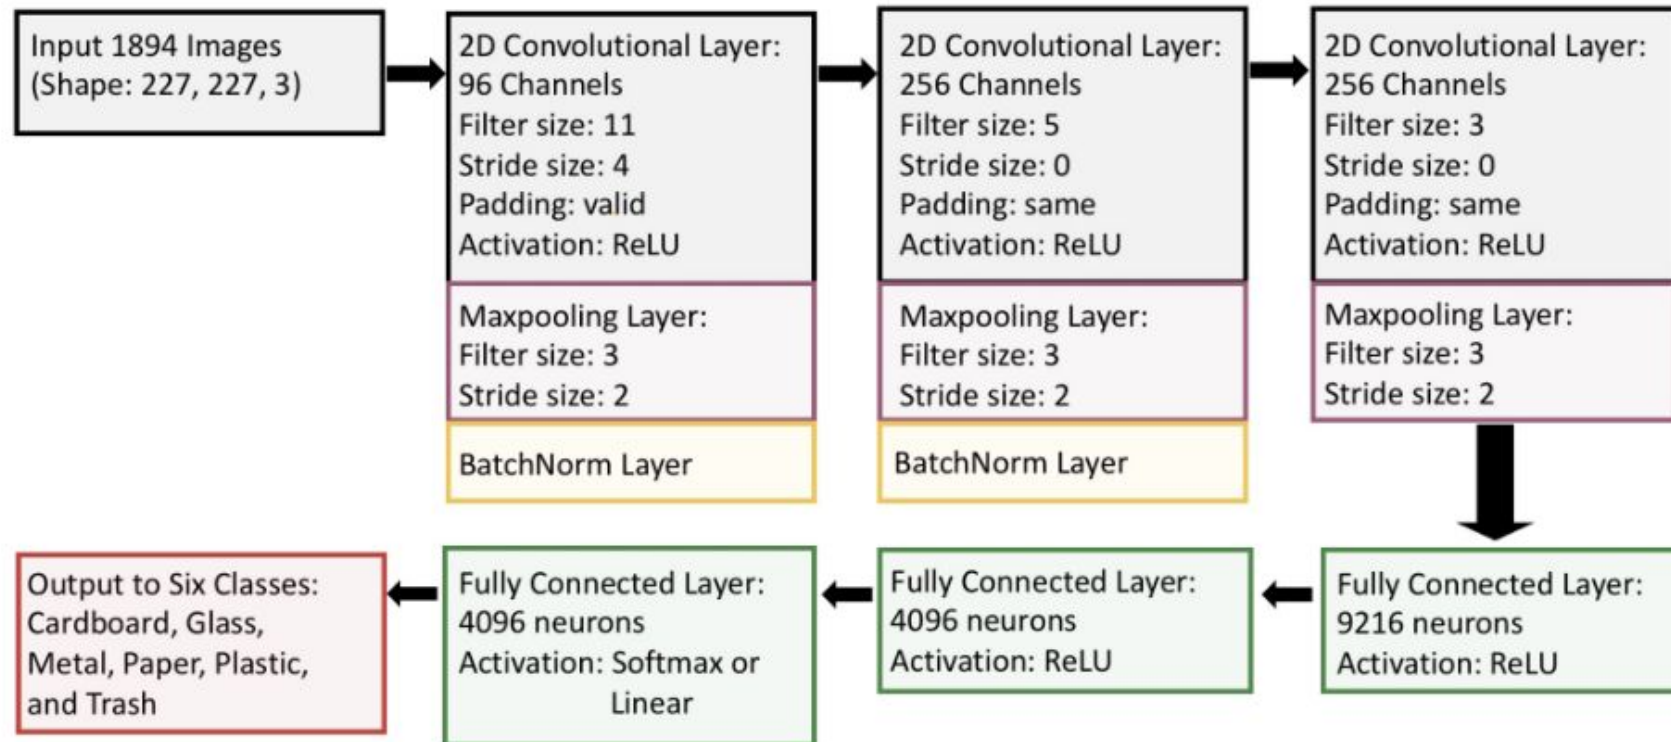




Figure 2: Basic Model Architecture







Baseline Outputs

Accuracy: 0.8709

Loss: 0.3401

Blue indicates what the model is focusing on



Adding Dropout

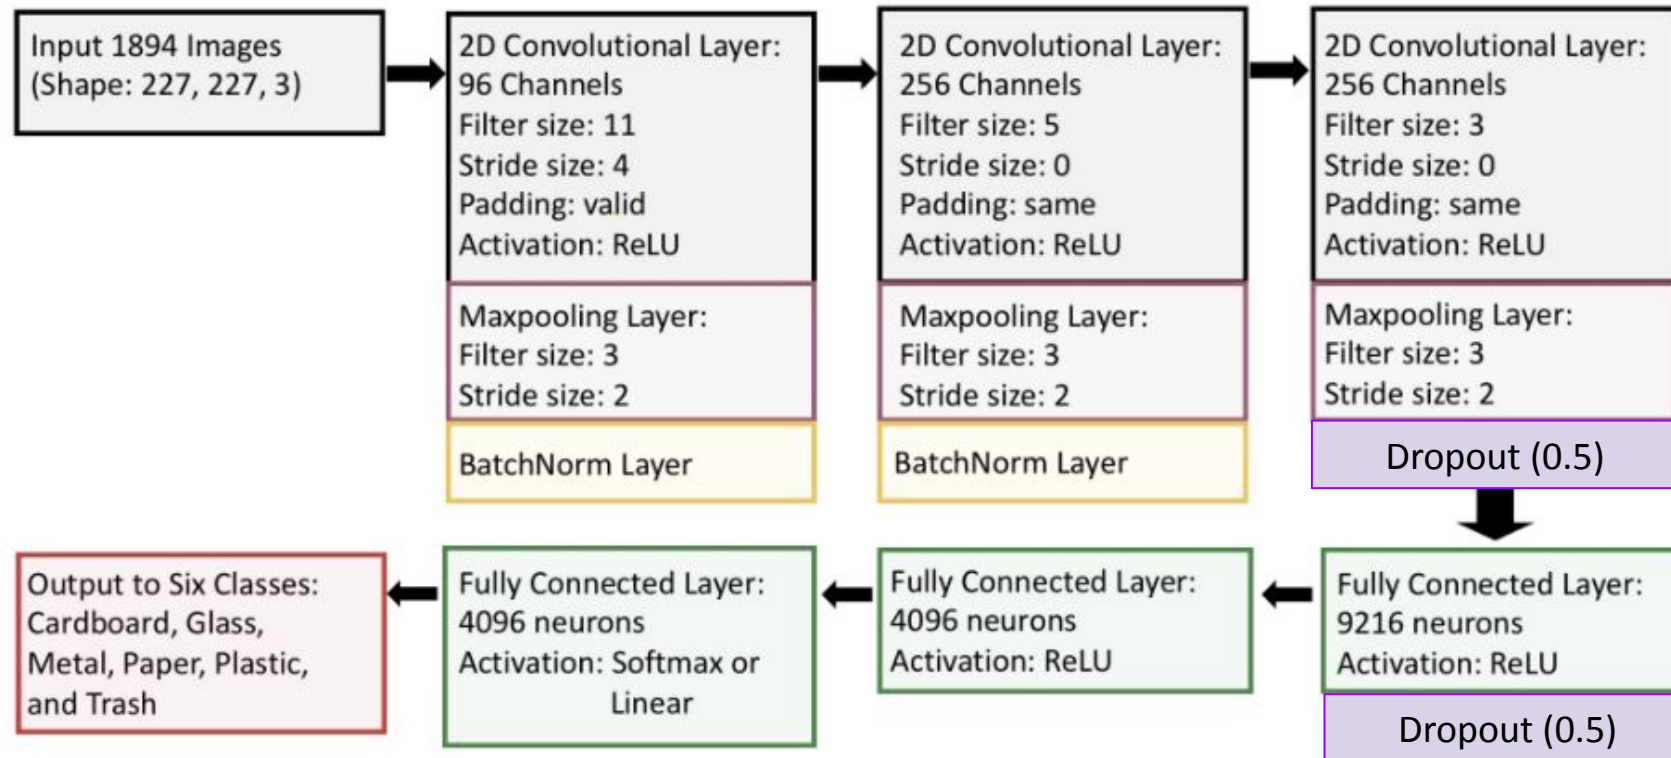




Figure 2: Basic Model Architecture





Dropout Change

Accuracy: 0.6147 (0.8709)

Loss: 0.9988 (0.3401)

Blue indicates what the model is focusing on



Removing Normalization

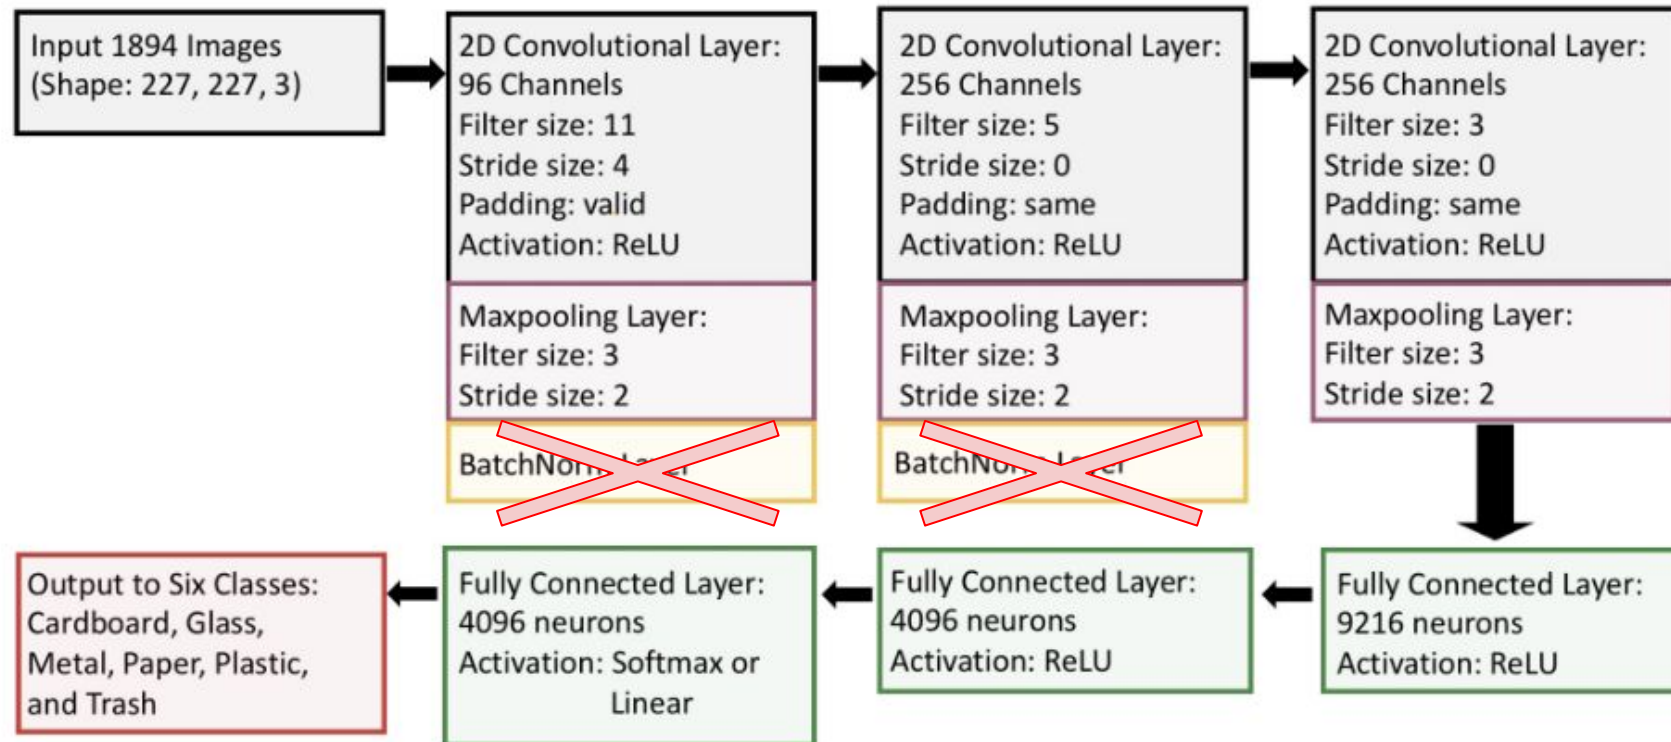




Figure 2: Basic Model Architecture




Normalization Change

Accuracy: 0.8161 (0.8709)

Loss: 0.4870 (0.3401)

Blue indicates what the model is focusing on



Removing ReLU

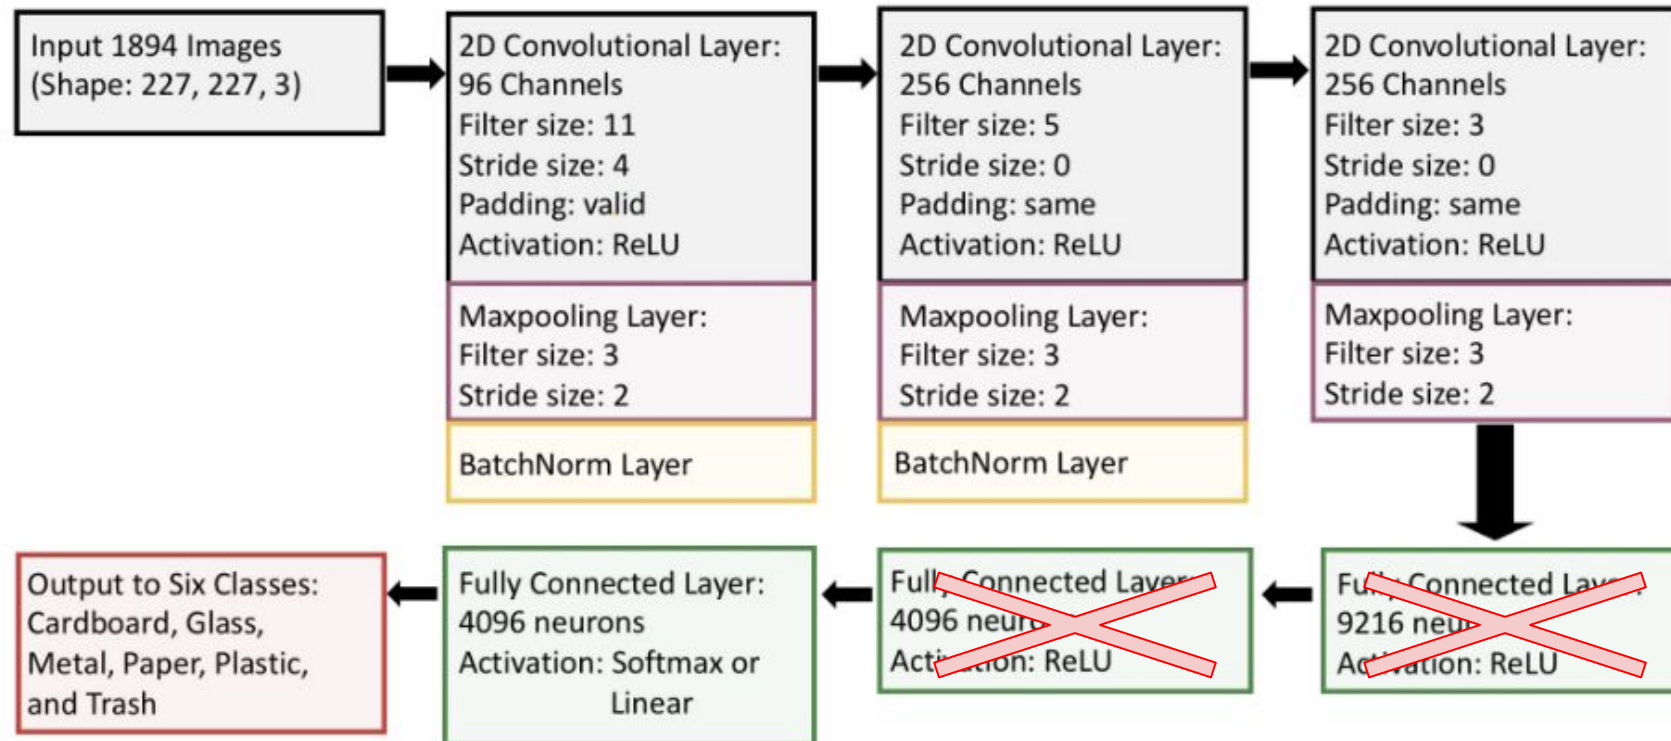







Figure 2: Basic Model Architecture



ReLU Change

Accuracy: 0.6021 (0.8709)

Loss: 1.1777 (0.3401)



Blue indicates what the model is focusing on

Removing CNN layers

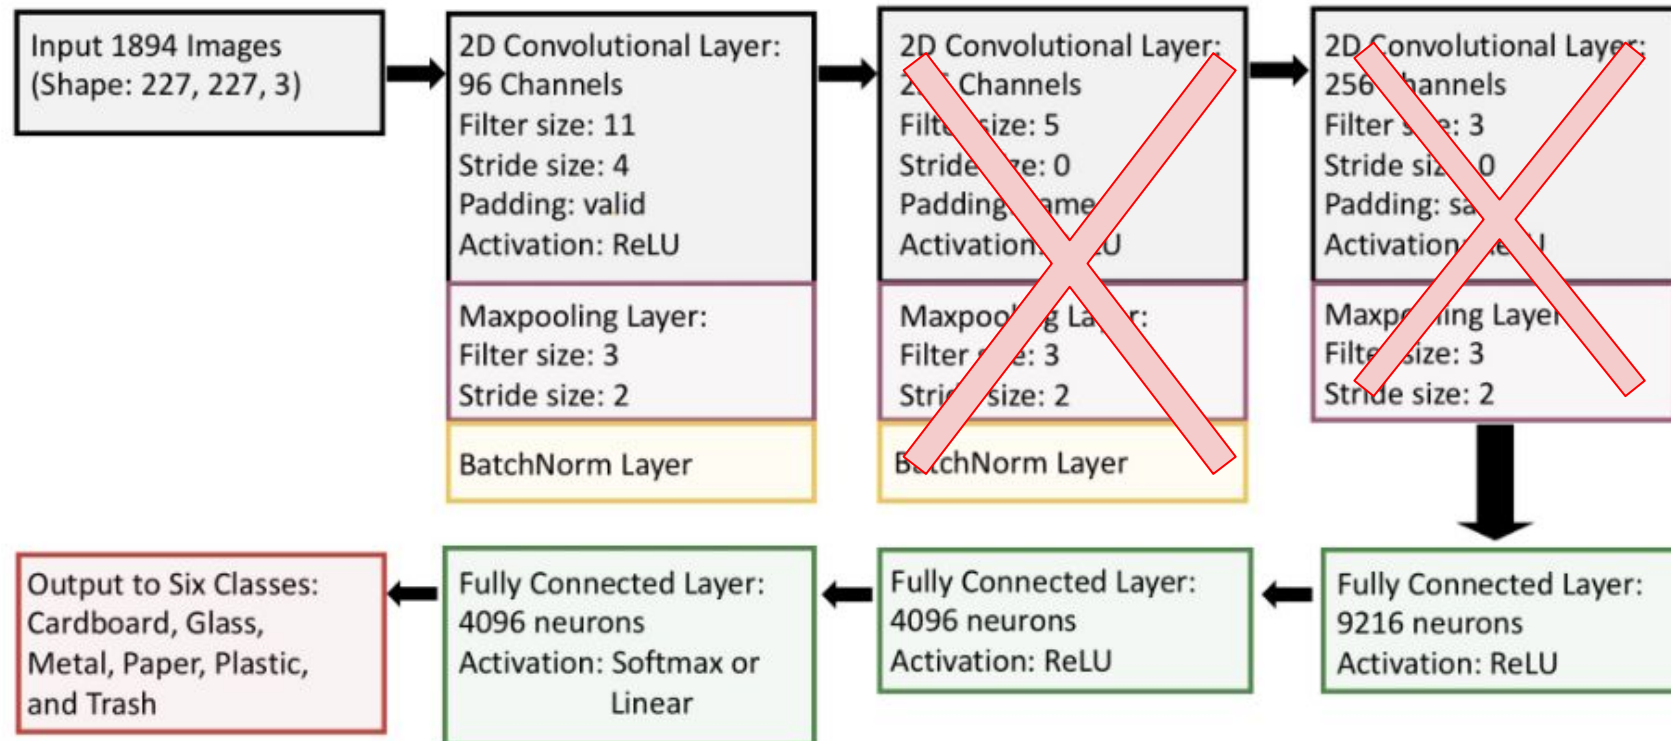






Figure 2: Basic Model Architecture



CNN Changes

Accuracy: 0.8920 (0.8709)

Loss: 0.2812 (0.3401)



Blue indicates what the model is focusing on

Notable Differences



Baseline




Added Dropout

Accuracy: B: 0.8709 N: 0.6147

Loss: B: 0.3401 N: 0.9988

Baseline model focuses better and has a better accuracy, shows dropout is not necessary for this trash classification.



Baseline




No Normalization

Accuracy: B: 0.8709 N: 0.8161

Loss: B: 0.3401 N: 0.4870

While accuracy and loss are similar for both models, explainable AI clearly shows that normalization is needed, as it doesn't focus on the can when removed.

Notable Differences



Baseline




No ReLU

Accuracy: B: 0.8709 N: 0.6021

Loss: B: 0.3401 N: 1.1777

ReLU seems to help focus the model and make the focus more intense. It should be kept in the model.



Baseline



No CNN

Accuracy: B: 0.8709 N: 0.8920

Loss: B: 0.3401 N: 0.2812

Removing the 2 CNN layers makes the model lose focus and intensity. While the accuracy and loss get better, the model appears to have sparse attention.



Challenges

- Issues uploading Dataset
- Updating old code for image augmentation
 - Converting images to same dimensions (227 x 227), was 512 X 384
 - Dimension inflation issues with Grad-CAM
- Space and time issues
 - Extremely long runtimes which limited our scope
 - limited GPU access



Conclusion:

Our findings suggest that they chose the right elements for their model as the baseline was the best in accuracy, loss, and where the model was focusing.

For the future

We ended up implementing 2 chunks of code in order to get all of these comparisons. We would recommend utilizing this tool whenever you decide to improve a model (given sufficient resources).

For Milestone 2, we would hope to have the resources to

- compare AlexNet with ResNet and VGG.
- Implement saliency maps
- Further vertical comparisons, add more to the model
- Global explanation
- Incorporate Mask-R-CNN to differ between overlapped classes (plastic)

```
grad_cam(model, image, class_idx):
    """Generates a heat map using Grad-CAM."""
    print(f"Inside grad_cam, image shape: {image.shape}")

    # Expand dimensions of the image to match input shape
    image = np.expand_dims(image, axis=0) # Ensure batch dimension is included

    # Get the gradient model
    grad_model = tf.keras.Model(
        inputs=model.inputs, outputs=model.get_layer("conv5").output, model.output
    )

    with tf.GradientTape() as tape:
        last_conv_layer = model.get_layer('conv5') # Ensure this matches your model's last conv layer
        iterate = tf.keras.models.Model([model.input], [last_conv_layer.output, model.output])

        # Use the already expanded image
        conv_outputs, preds = iterate(image) # Image is now (1, 227, 227, 3)
        loss = preds[0, class_idx]

    grads = tape.gradient(loss, conv_outputs)[0]
    pooled_grads = tf.reduce_mean(grads, axis=(0, 1)) # Average over the spatial dimensions

    # Convert the conv_outputs to NumPy array
    heatmap = conv_outputs[0].numpy() # Convert to NumPy array for manipulation

    print(f"Shape of pooled_grads: {pooled_grads.shape}") # Debugging line
    print(f"Shape of heatmap before scaling: {heatmap.shape}") # Debugging line

    # Ensure pooled_grads is a numpy array
    pooled_grads = pooled_grads.numpy() # Convert to NumPy for manipulation
    pooled_grads = np.expand_dims(pooled_grads, axis=(0, 1)) # Shape becomes (1, 1, 256)

    # Check shapes before multiplication
    print(f"Shape of pooled_grads after expansion: {pooled_grads.shape}") # Debugging line
    print(f"Shape of heatmap before applying pooled_grads: {heatmap.shape}") # Debugging line

    # Multiply each channel in the heatmap with the corresponding pooled gradient
    # Ensure correct dimensions for multiplication
    heatmap = np.tensordot(heatmap, pooled_grads, axes=(2, 2)) # Shape will be (13, 13)

    heatmap = np.maximum(heatmap, 0) # Ensure non-negative values
    heatmap /= np.max(heatmap) # Normalize to [0, 1]

    # Squeeze the heatmap
    heatmap = np.squeeze(heatmap) # Remove dimensions of size 1

    # Check heatmap shape before resizing
    print(f"Heatmap shape before resizing: {heatmap.shape}")

    # Resize heatmap to match the input image dimensions
    if heatmap.ndim == 2: # Ensure the heatmap is 2D
        heatmap = cv2.resize(heatmap, (image.shape[2], image.shape[3]))
    else:
        raise ValueError("Heatmap does not have the expected dimensions")

    return heatmap
```

```
display_grad_cam(image, heatmap):
    # Resize heatmap to match the original image size
    heatmap = np.uint8(255 * heatmap) # Scale to 0-255
    heatmap = cv2.applyColorMap(heatmap, cv2.COLORMAP_JET) # Apply color map

    # Overlay the heatmap on the original image
    superimposed_img = heatmap * 0.4 + image # Adjust alpha as needed
    plt.imshow(superimposed_img / 255.0) # Display as float image
    plt.axis('off')
    plt.show()

# Run Grad-CAM on the selected images
print(f"Number of selected images: {len(selected_images)}") # Debug statement

for idx, selected_image in enumerate(selected_images):
    print(f"Selected image {idx} shape: {selected_image.shape}") # Debug statement

    # Check if the loop is entering the body
    if selected_image is None:
        print(f"Image {idx} is None.") # Check for None
        continue # Skip if None

    # Ensure we expand the dimensions correctly for the model
    expanded_image = np.expand_dims(selected_image, axis=0)
    print(f"After expansion, expanded image shape: {expanded_image.shape}")

    # Predict class using the model
    class_idx = np.argmax(model.predict(expanded_image))
    print(f"Image {idx}: Predicted class index: {class_idx}")

    # Call Grad-CAM
    print(f"Calling grad_cam for image {idx} with class index {class_idx}")
    heatmap = grad_cam(model, np.squeeze(selected_image), class_idx) # Remove extra dimensions
    display_grad_cam(selected_image, heatmap)
```



Q&A